



Cloud Native ML with Kubeflow & TensorRT



A network diagram showing a central node with 'mem' and 'proc' labels, connected to several other nodes. Below it, there are labels 'P2', 'P1', 'P2', and 'A1' arranged vertically. A horizontal bar with a gradient is also visible.

David Aronchick
Head of OSS ML Strategy, Azure

Seth Juarez
Senior Cloud Developer Advocate, Azure



One Year Ago...

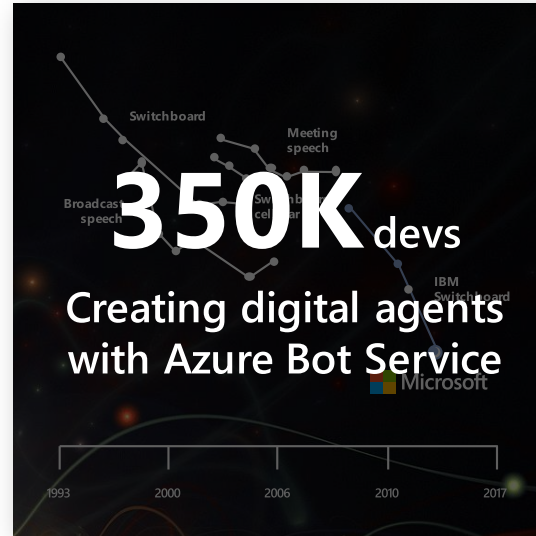
What is Machine Learning?

**Machine Learning is a
way of solving problems
without explicitly
knowing how to create
the solution.**



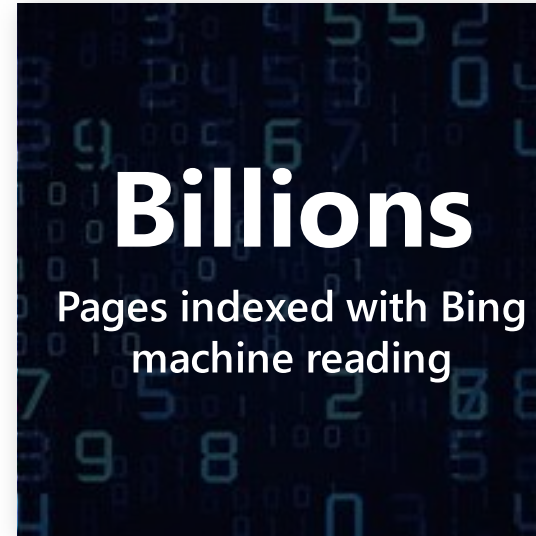
2016

Object recognition
Human parity



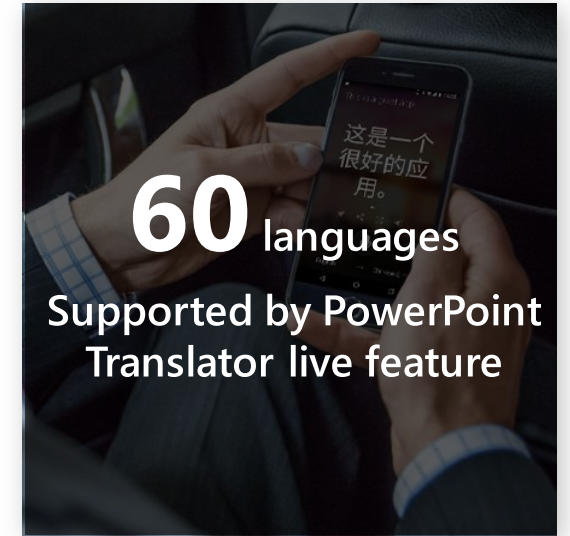
2017

Speech recognition
Human parity



January 2018

Machine reading comprehension
Human parity



March 2018

Machine translation
Human parity



But ML is hard!

Four Years Ago...

Google and Containers

Everything at Google runs in a container.

Internal usage:

- Resource isolation and predictability
- Quality of Services
 - batch vs. latency sensitive serving
- Overcommitment (not for GCE)
- Resource Accounting

We start over 2 billion containers per week.





Kubernetes

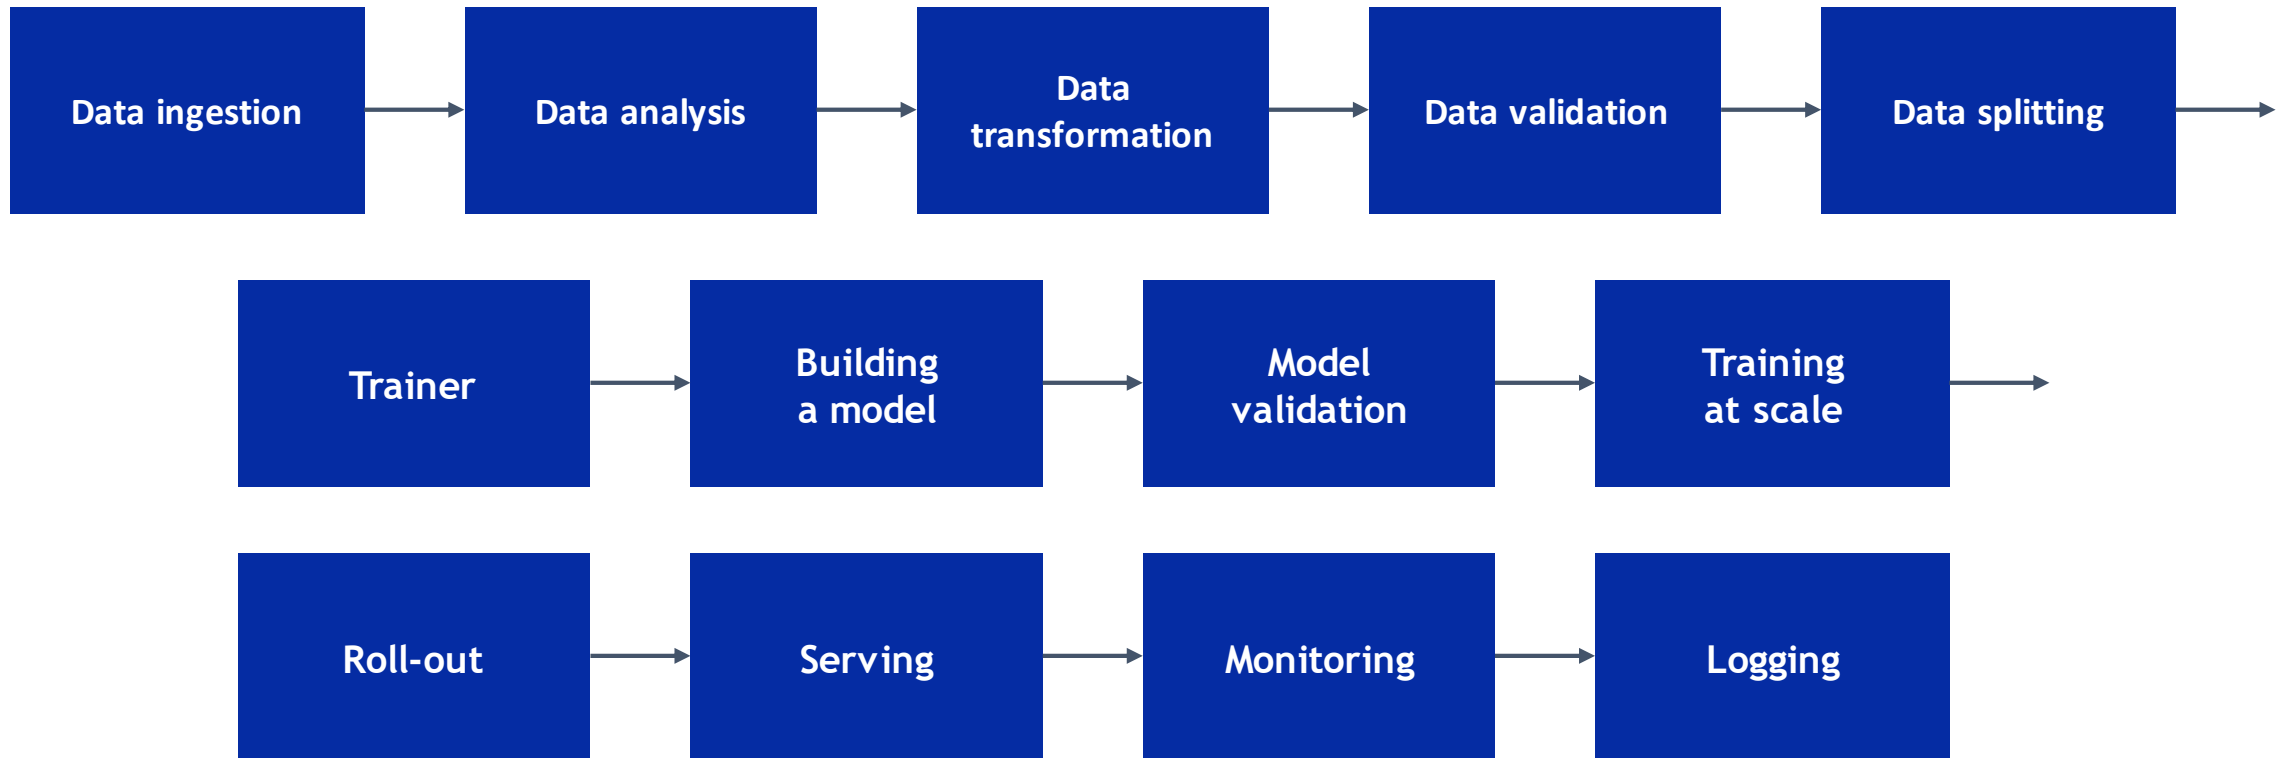
Cloud Native Apps

Cloud Native ML?

Platform

Building a model

Platform



Kubecon 2017

Introducing Kubeflow



KubeCon



CloudNativeCon

North America 2017

**Make it Easy for Everyone
to **Develop, Deploy** and **Manage**
Portable, Distributed ML
on Kubernetes**

Experimentation

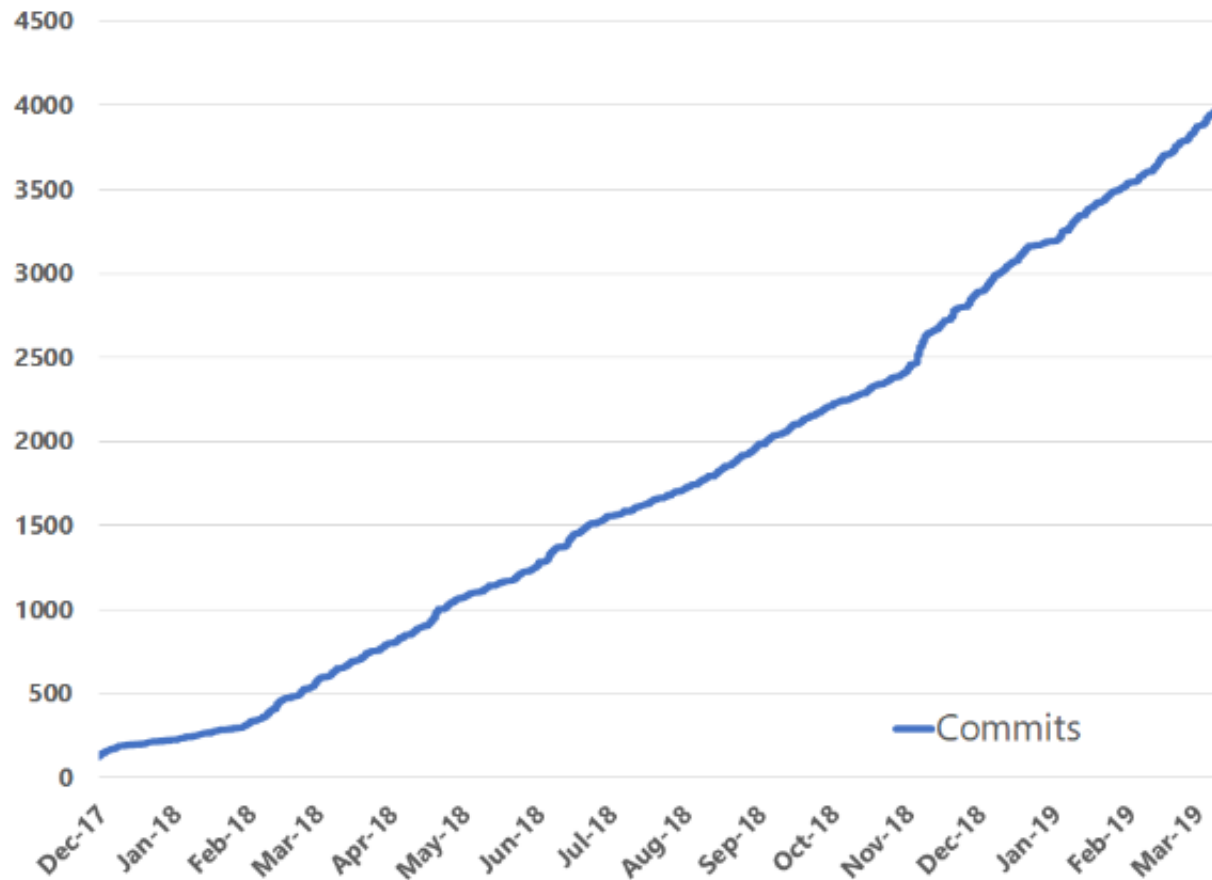
Training

Cloud



Cloud Native ML!

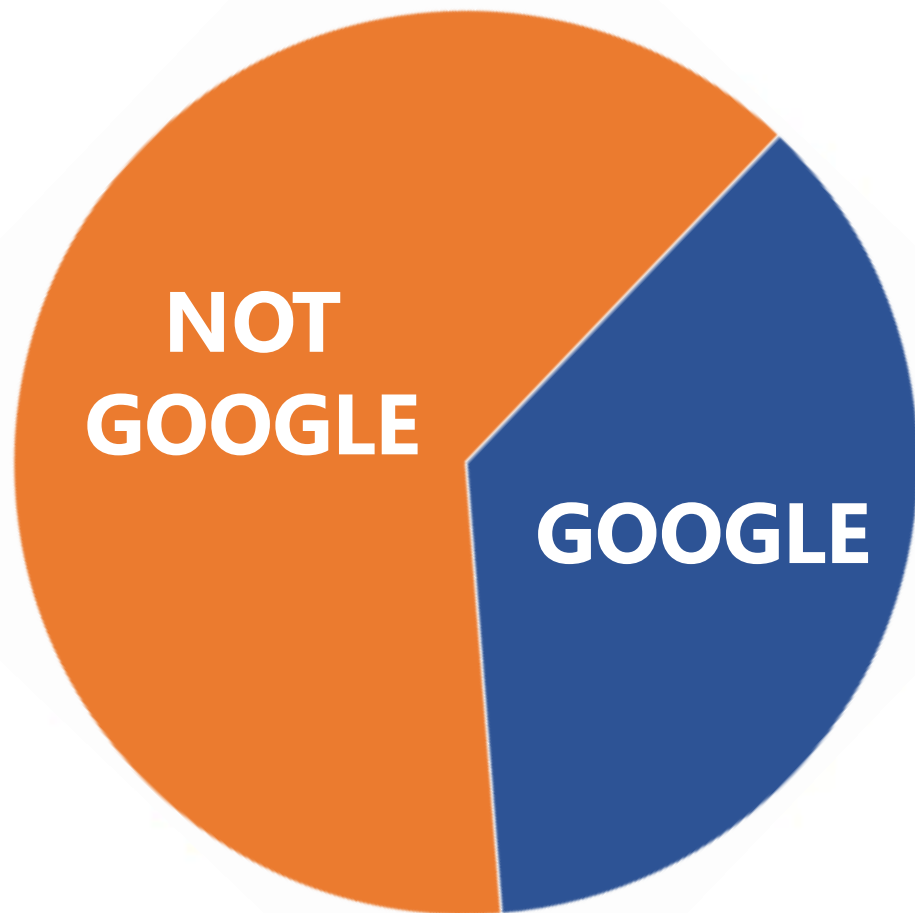
Momentum!



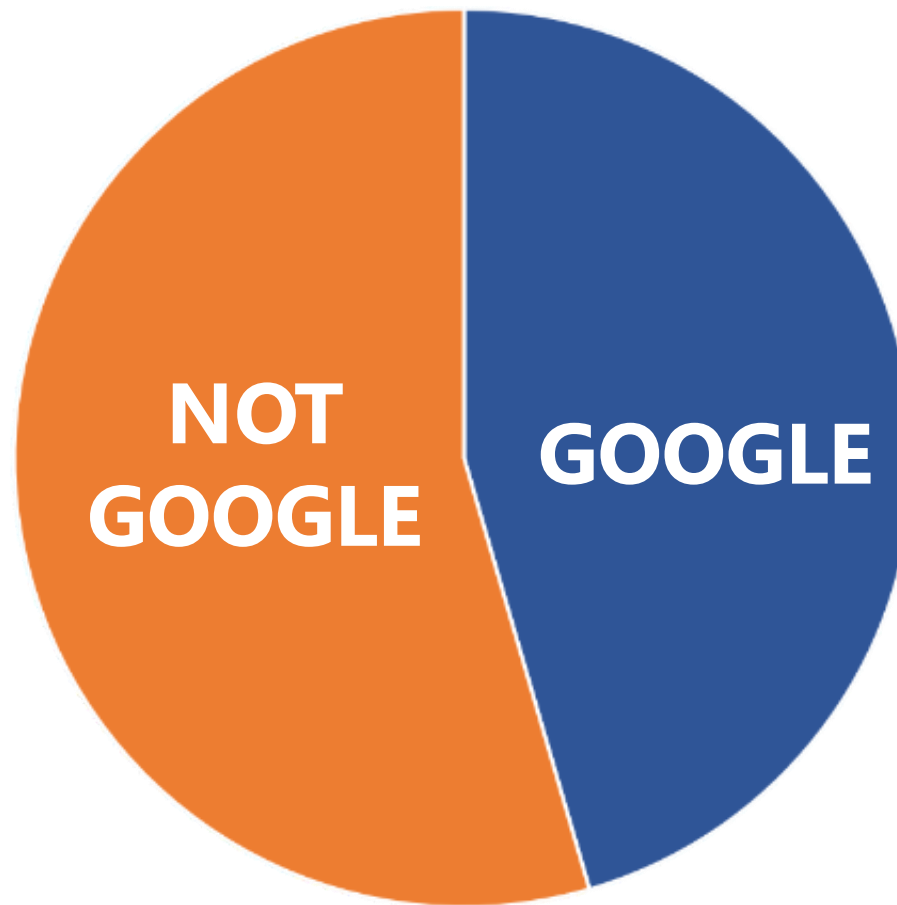
- ~4000 commits
- ~200 community contributors
- ~50 companies contributing, including:



Community Contributions



Kubernetes



Kubeflow

Critical User Journey Comparison

2017

- Experiment with Jupyter
- Distribute your training with TFJob
- Serve your model with TF Serving

2019

- Setup locally with miniKF
- Access your cluster with Istio/Ingress
- Ingest your data with Pachyderm
- Transform your data with TF.T
- Analyze the data with TF.DV
- Experiment with Jupyter
- Hyperparam sweep with Katib
- Distribute your training with TFJob
- Analyze your model with TF.MA
- Serve your model with Seldon
- Orchestrate everything with KF.Pipelines

Community Contribution

- Pluggable microservice architecture for HP tuning
 - Different optimization algorithms
 - Different frameworks
- StudyJob (K8s CRD)
 - Hides complexity from user
 - No code needed to do HP tuning

Katib from NTT

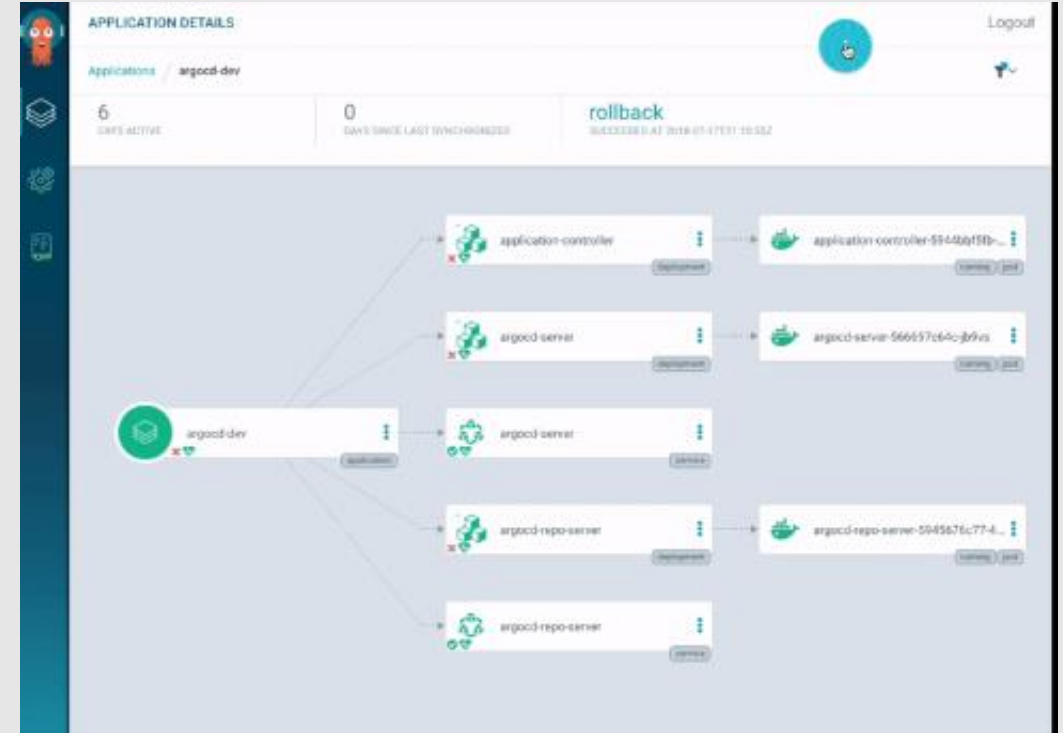
The screenshot shows the ModelDB web interface with three projects listed. Each project has a table of metrics and a link to Project Model Types.

Project	Validation Accuracy	Min Cost	Max Cost	Average
cifar10	0.94	0.94	0.94	0.94
cifar10-py-test	0.94	0.94	0.94	0.94
tf-resnet	0.94	0.94	0.94	0.94

Community Contribution

- Argo CRD for workflows
- Argo CRD is engine for Pipelines
- Argo CD for GitOps

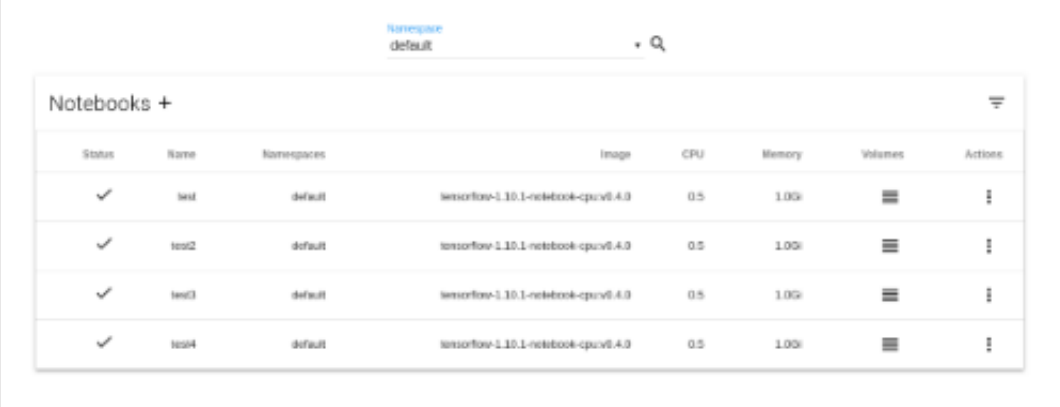
Argo from Intuit



Community Contribution

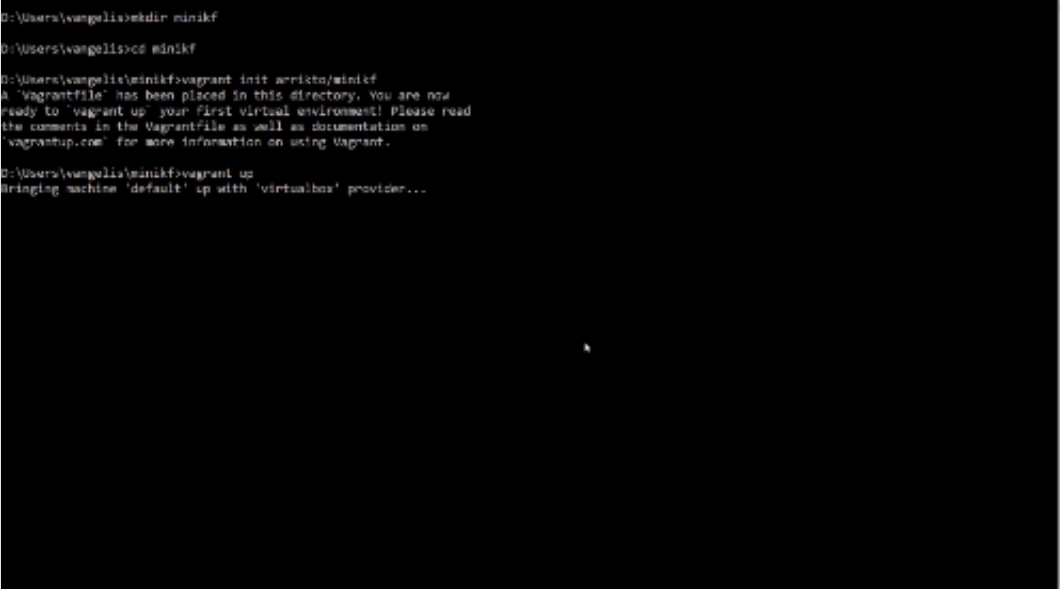
- Core Notebook Experience
 - **0.4:** New JupyterHub-based UI
 - **0.5:** K8s-Native Notebooks UI
- Pipelines: Support for local storage
- Multiple Persistent Volumes
- MiniKF: All-in-one packaging for seamless local deployments

NB & Storage from Arrikto



The screenshot shows a web interface for managing notebooks. At the top, there's a header with 'Namespace: default' and a search icon. Below this is a section titled 'Notebooks +' with a dropdown arrow. The main content is a table with columns: Status, Name, Namespace, Image, CPU, Memory, Volumes, and Actions. There are four rows of notebooks, all with a status of '✓' and namespace of 'default'. Each row has a 'Volumes' column with a hamburger menu icon and an 'Actions' column with a vertical ellipsis icon.

Status	Name	Namespace	Image	CPU	Memory	Volumes	Actions
✓	test	default	tensorflow-1.10.1-notebook-cpu-v6.4.0	0.5	1.0Gi	⋮	⋮
✓	test2	default	tensorflow-1.10.1-notebook-cpu-v6.4.0	0.5	1.0Gi	⋮	⋮
✓	test3	default	tensorflow-1.10.1-notebook-cpu-v6.4.0	0.5	1.0Gi	⋮	⋮
✓	test4	default	tensorflow-1.10.1-notebook-cpu-v6.4.0	0.5	1.0Gi	⋮	⋮



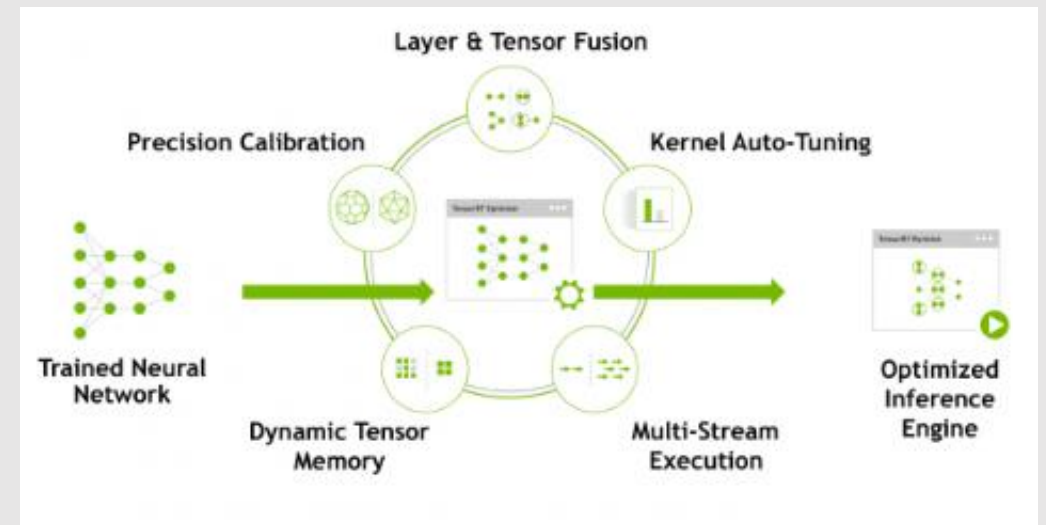
The terminal shows the following commands and output:

```
D:\Users\vangelis>cd minikf
D:\Users\vangelis>cd minikf
D:\Users\vangelis\minikf>vagrant init arrikto/minikf
A 'Vagrantfile' has been placed in this directory. You are now
ready to 'vagrant up' your first virtual environment! Please read
the comments in the Vagrantfile as well as documentation on
'vagrantup.com' for more information on using Vagrant.
D:\Users\vangelis\minikf>vagrant up
Bringing machine 'default' up with 'virtualbox' provider...
```

Community Contribution

- Production datacenter inferencing server
- Maximize real-time inference performance of GPUs
- Multiple models per GPU per node
- Supports heterogeneous GPUs & multi GPU nodes
- Integrates with orchestration systems and auto scalers via latency and health metrics

TensorRT from NVidia



Introducing Kubeflow 0.5

What's in the box?

UX investments - First class notebooks & central dashboard

- Build/Train/Deploy From notebook
- Better multi-user support
- A new web-based spawer

Enterprise readiness

- Better namespace support
- API stability
- Upgradability with preservation of historical metadata

Advanced composability & tooling

- Advanced support for calling out to web services
- Ability to specify GPU/TPUs for pipeline steps
- New metadata backend



Better/Faster/Production Notebooks!

User Goal = Just give me a notebook!

Problem

- Setting up A notebook is $O(\text{easy})$
- Setting up a rich, production-ready notebook is $O(\text{hard})$
- Setting up a rich, production-ready notebook that works anywhere, on any cloud, with a minimum of changes is $O(\text{very very hard})$

Better/Faster/Production Notebooks!

Setting up a notebook is easy!

```
$ curl -O
https://repo.continuum.io/archive/Anaconda3-5.0.1-
Linux-x86_64.sh

$ bash -c Anaconda3-5.0.1-Linux-x86_64.sh

$ conda create -y -n mlenv python=2 pip scipy
gevent sympy

$ source activate mlenv

$ pip install tensorflow==1.13.0 | tensorflow-
gpu==1.7.0

$ open http://127.0.0.1:8080
```



Except...

- Custom libraries
- HW provisioning (especially GPUs) & **drivers**
- Portability (between laptop and clouds)
- Security profiles
- Service accounts
- Credentials
- Lots more...



Better/Faster/Production Notebooks!

Solution –

**Declarative Data Science Environments with
Kubeflow!**

Better/Faster/Production Notebooks!

Setting up a declarative environment is easy!

```
$ kfctl.sh init
$ kfctl.sh --platform aks \
            --project my-project
$ kfctl.sh generate platform
$ kfctl.sh apply platform
$ kfctl.sh generate k8s
$ kfctl.sh apply k8s
```

Add your custom components!

```
# Add Seldon Server
$ ks pkg install kubeflow/seldon

# Add XGBoost
$ ks pkg install kubeflow/xgboost

# Add hyperparameter tuning
$ ks pkg install kubeflow/katib

# Add Seldon Server
$ ks pkg install kubeflow/seldon
```

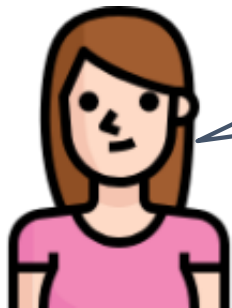
Experimentation

Training

Cloud



IT
Ops



I Got You!



DEMO

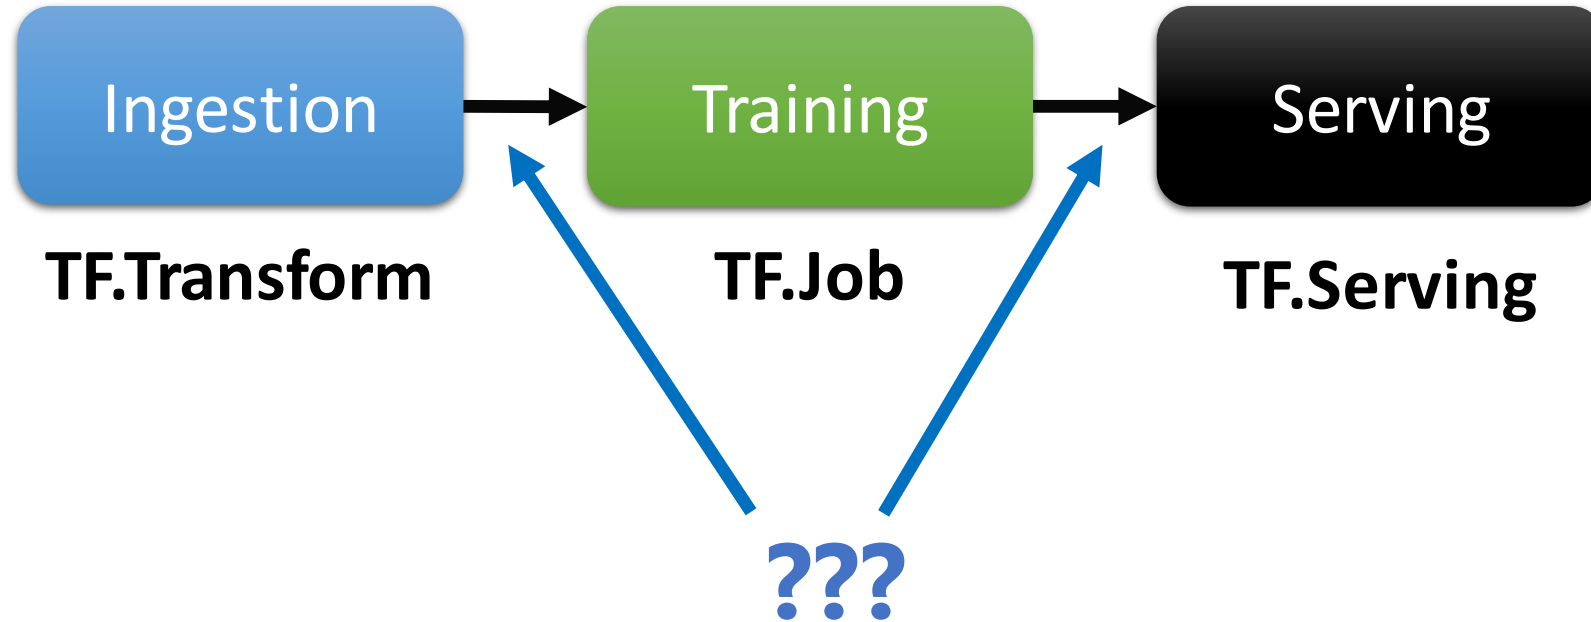
Rich Container Based Pipelines

User Goal = Repeatable, multi-stage ML training

Problem

- Tools not built to be containerized/orchestrated
- Coordinating between steps often requires writing custom code
- Different tools have different infra requirements

Rich Container Based Pipelines



Pipelines should:

- Be cloud native (microservice oriented, loosely coupled) and ML aware
- Support both data and task driven workflows
- Understand non-Kubeflow-based services (e.g. external to the cluster)

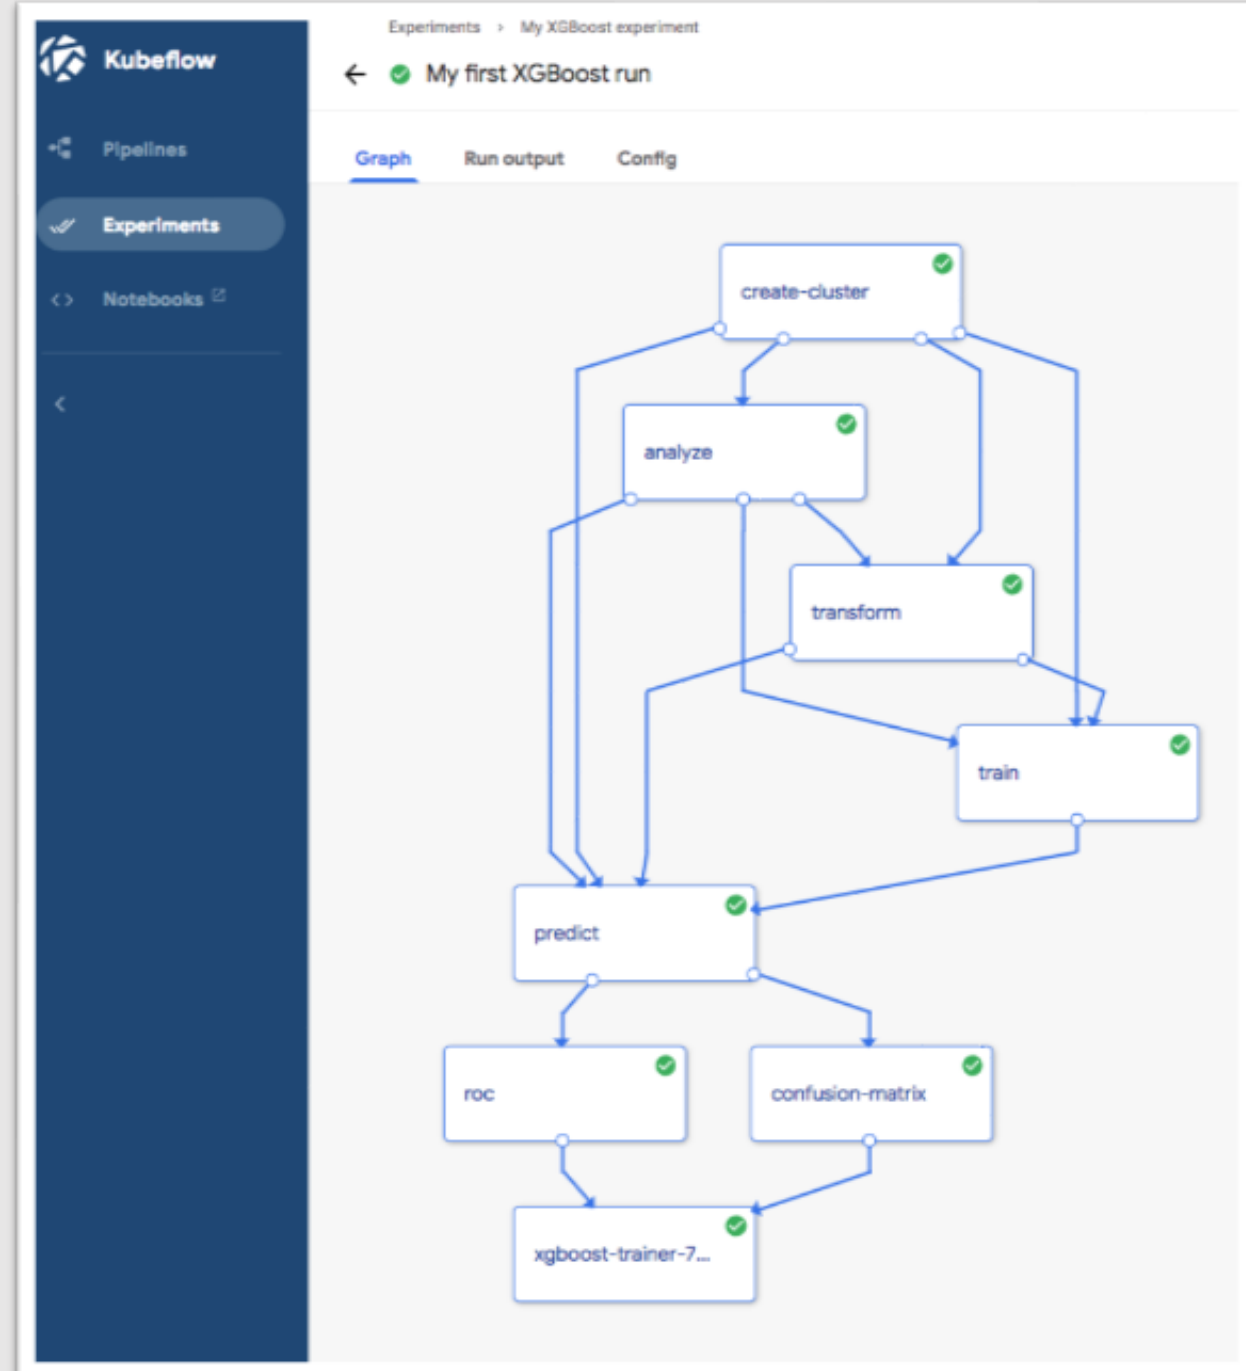
Rich Container Based Pipelines

Solution –

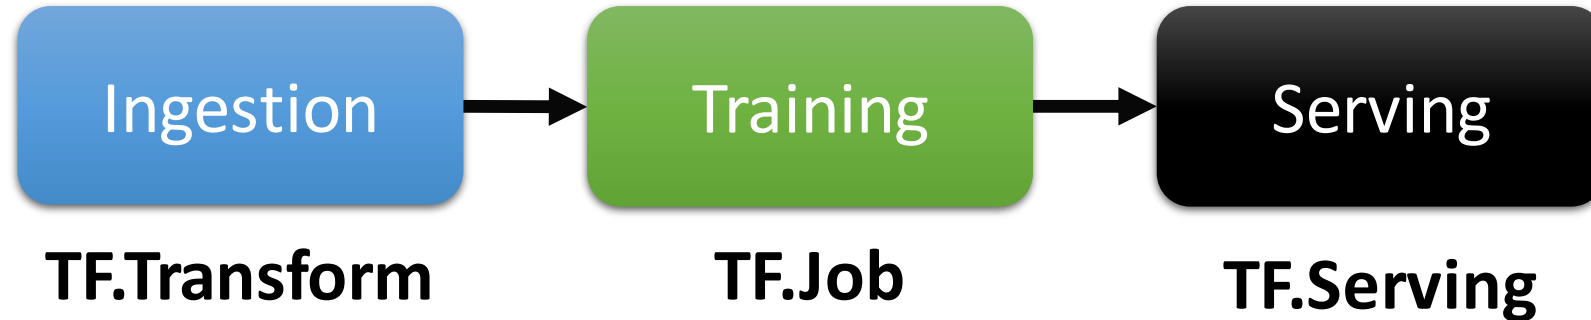
Kubeflow Pipelines!

Kubeflow Pipeline Details

- Containerized Implementations of ML Tasks
 - Escapsulates all the dependencies of a step with no conflicts
 - Step can be singular or distributed
 - Can also involve external services
- Specified via Python SDK
- Inputs/outputs/parameters can be chained together



Rich Container Based Pipelines



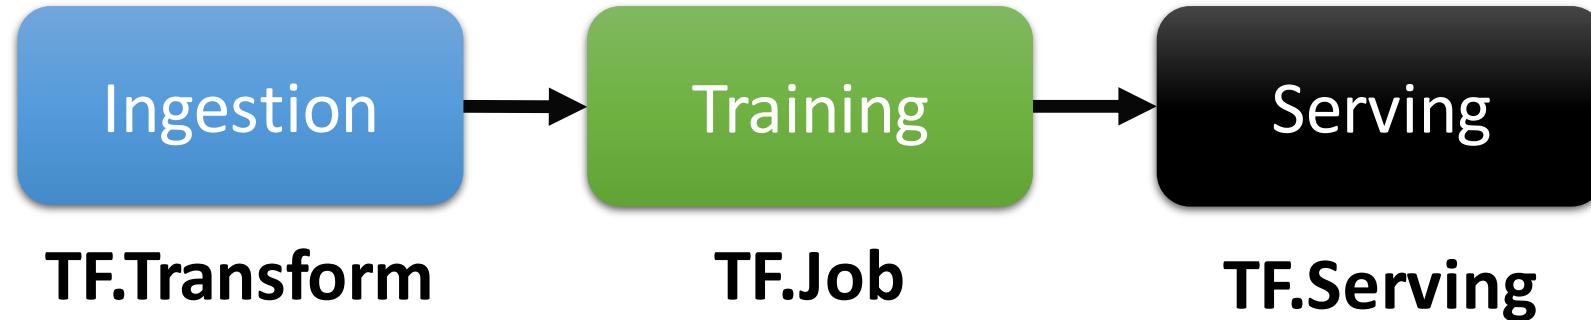
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=tfs_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

**Can I Change
a Step?**

Rich Container Based Pipelines



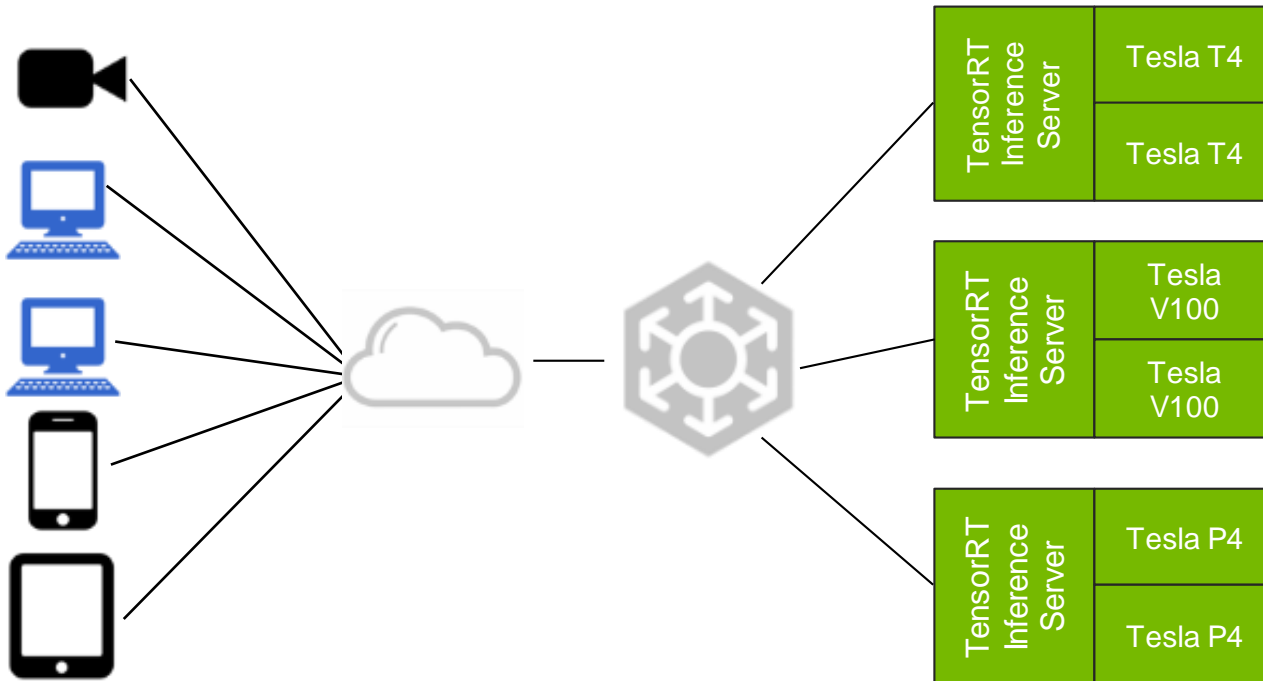
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=tfs_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

NVIDIA TENSORRT INFERENCE SERVER

Production Data Center Inference Server



Maximize inference throughput & GPU utilization

Quickly deploy and manage multiple models per GPU per node

Easily scale to heterogeneous GPUs and multi GPU nodes

Integrates with orchestration systems and auto scalers via latency and health metrics

Now open source for thorough customization and integration

FEATURES

Concurrent Model Execution

Multiple models (or multiple instances of same model) may execute on GPU simultaneously

Eager Model Loading

Any mix of models specified at server start. All models loaded into memory.

CPU Model Inference Execution

Framework native models can execute inference requests on the CPU

Metrics

Utilization, count, and latency

Custom Backend

Custom backend allows the user more flexibility by providing their own implementation of an execution engine through the use of a shared library

Dynamic Batching

Inference requests can be batched up by the inference server to 1) the model-allowed maximum or 2) the user-defined latency SLA

Multiple Model Format Support

TensorFlow GraphDef / SavedModel
TensorFlow and TensorRT GraphDef
TensorRT Plans
Caffe2 NetDef (ONNX import path)

Mounted Model Repository

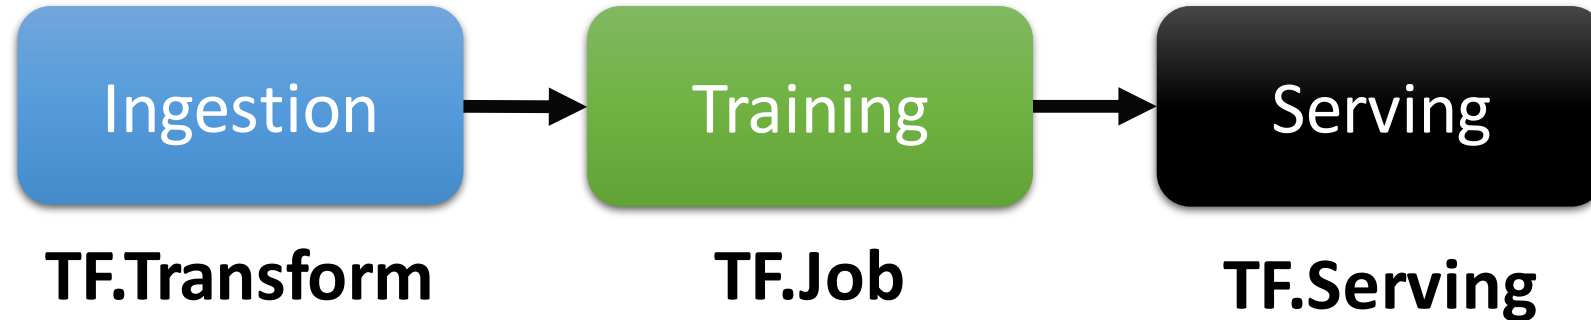
Models must be stored on a locally accessible mount point



TensorRT



Rich Container Based Pipelines

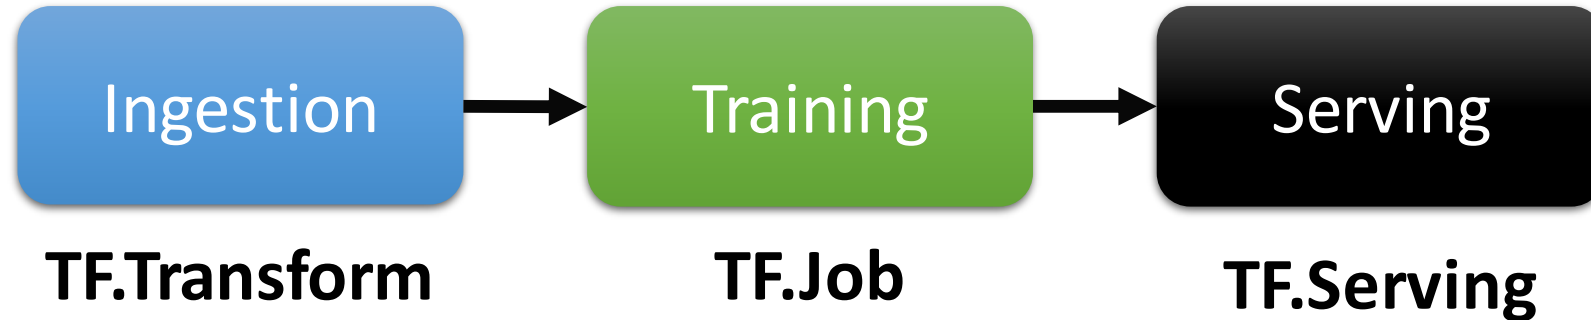


```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=tfs_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```


Rich Container Based Pipelines

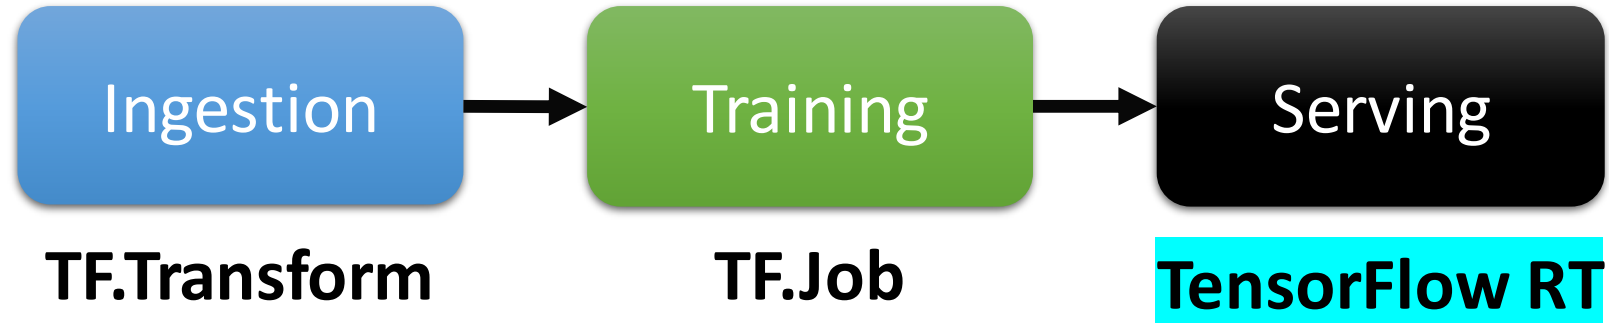


```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



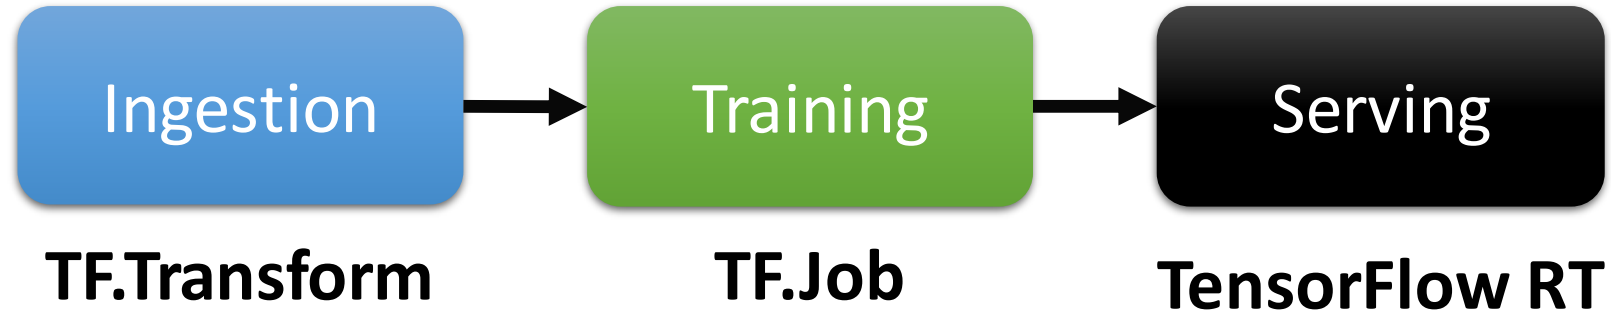
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

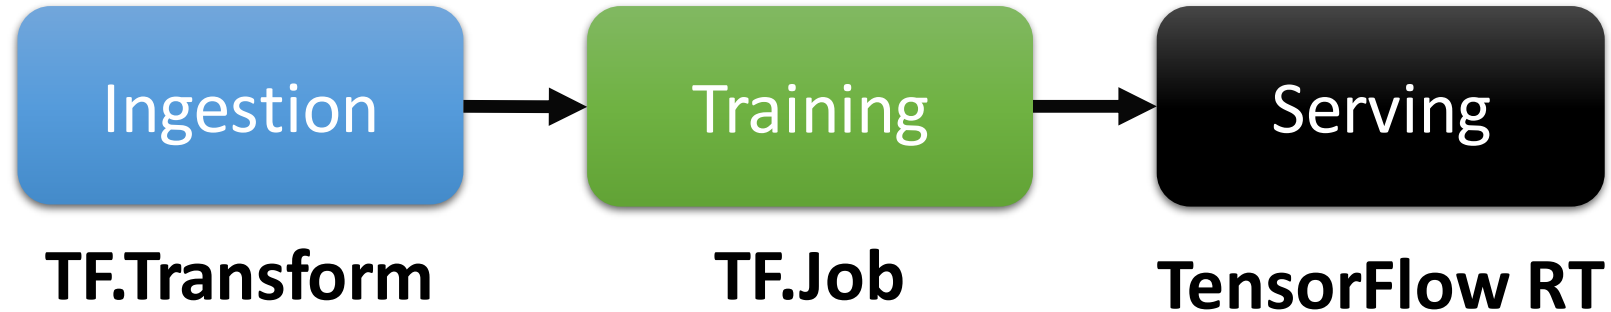
Now, Add a Step

Rich Container Based Pipelines



```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})  
  
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])  
  
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



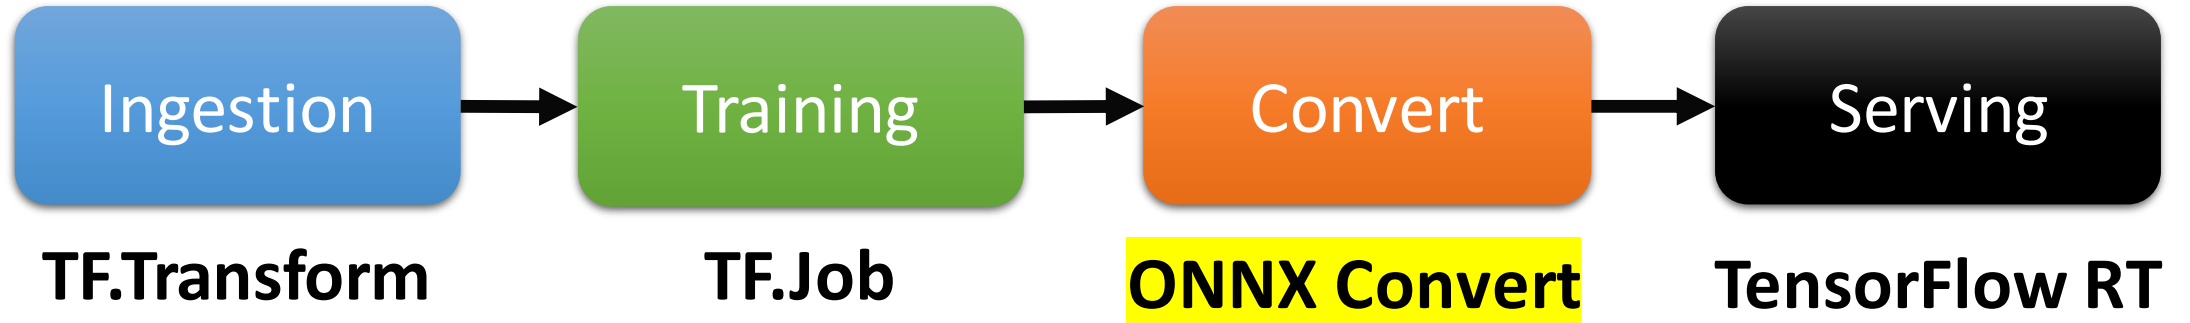
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



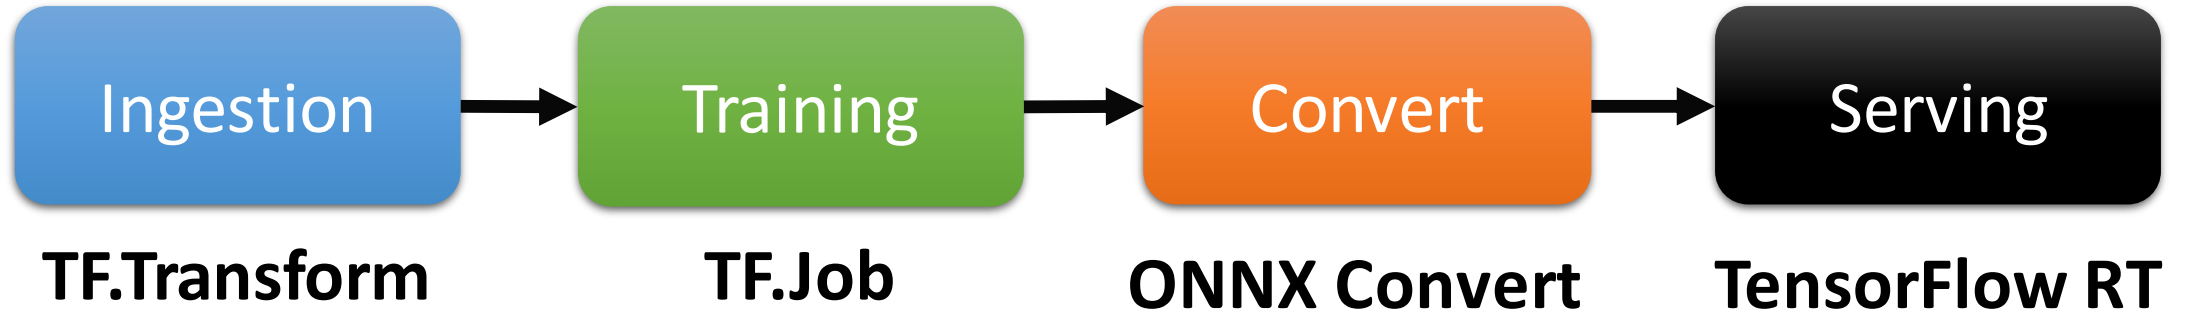
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Kubeflow Pipeline 0.5

- **UI/UX/SDK improvements**

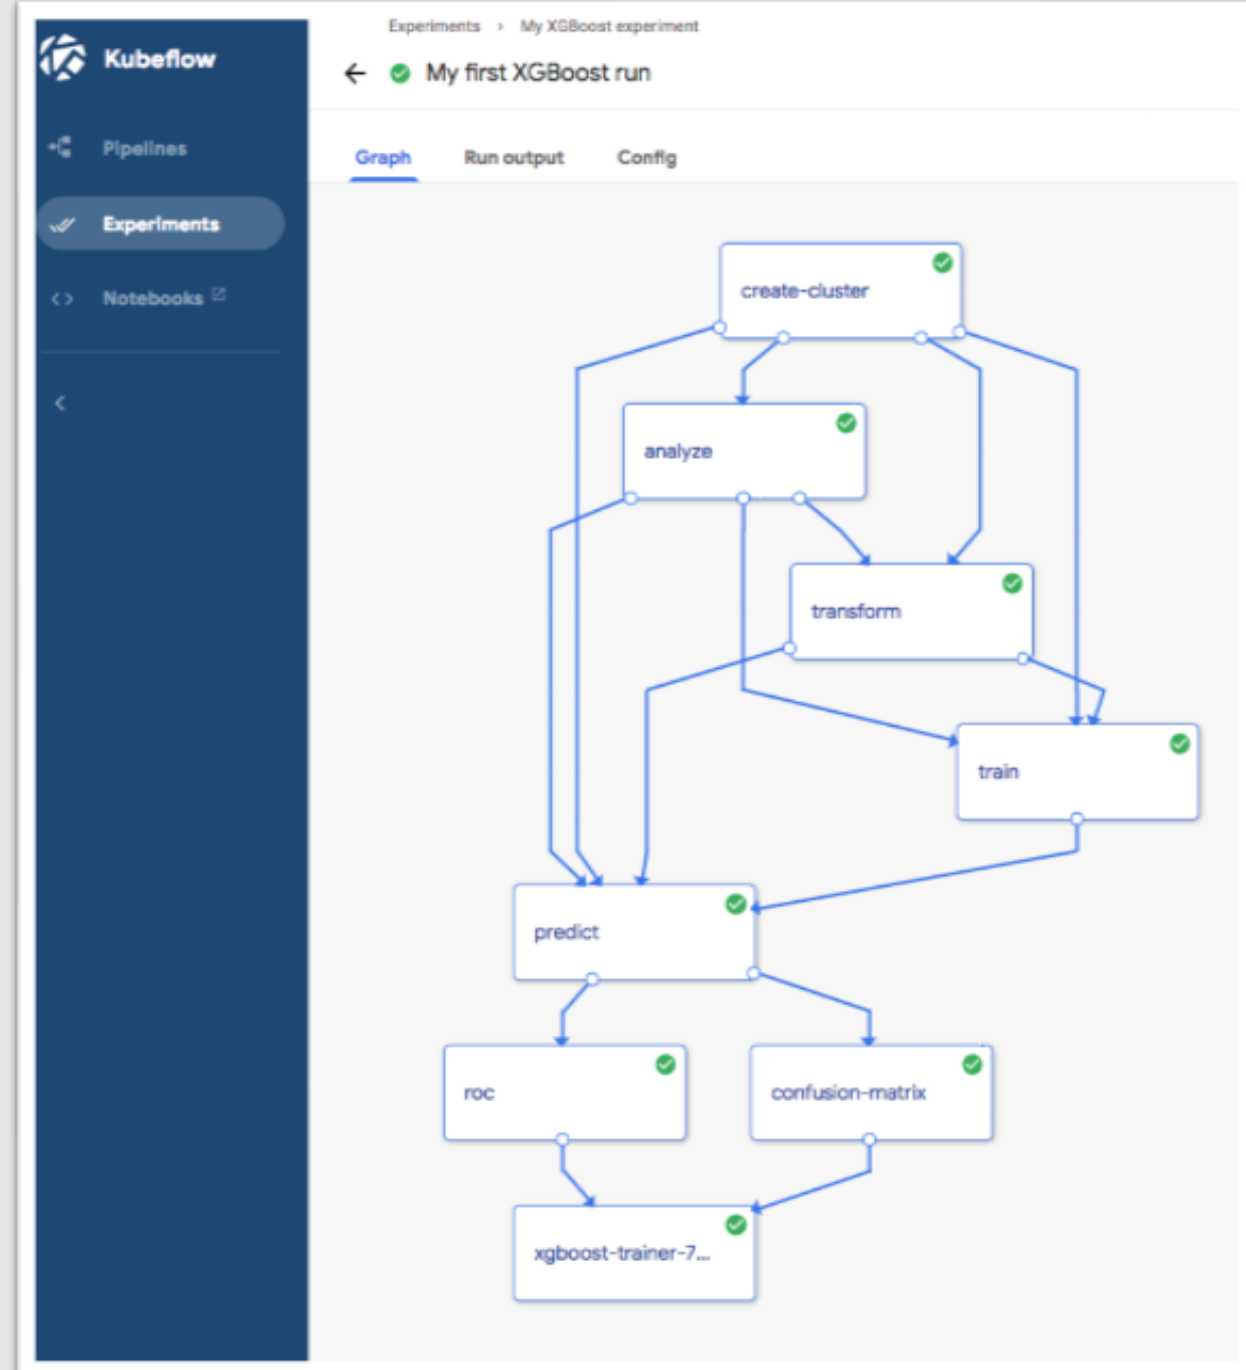
- Ability to specify GPU/TPU for pipeline steps
- Improved job search
- Metadata Backend to store and query metadata about artifacts produced by pipeline steps

- **Production readiness**

- Ability to upgrade a cluster without losing information about past runs
- Lots of stability improvements

- **Improved composability**

- Define and easily re-use a pipeline component.
- Compose a larger pipeline using smaller pipelines as building blocks



DEMO

Integrate External Services into Pipeline

User Goal = Just deploy manage it for me

Problem

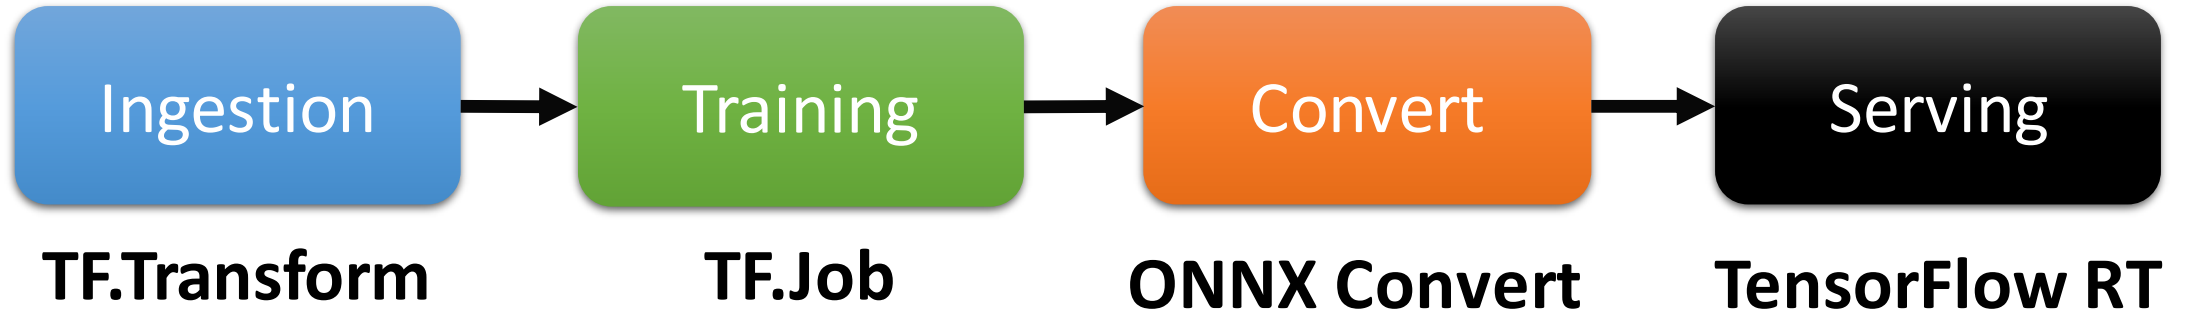
- Self-hosting is customizable but requires (too much) management
- Production requirements for model hosting
- Ability to scale dynamically based on demand, uptime, etc.

Integrate External Services into Pipeline

Solution –

**Kubeflow Pipelines!
(Again!)**

Rich Container Based Pipelines



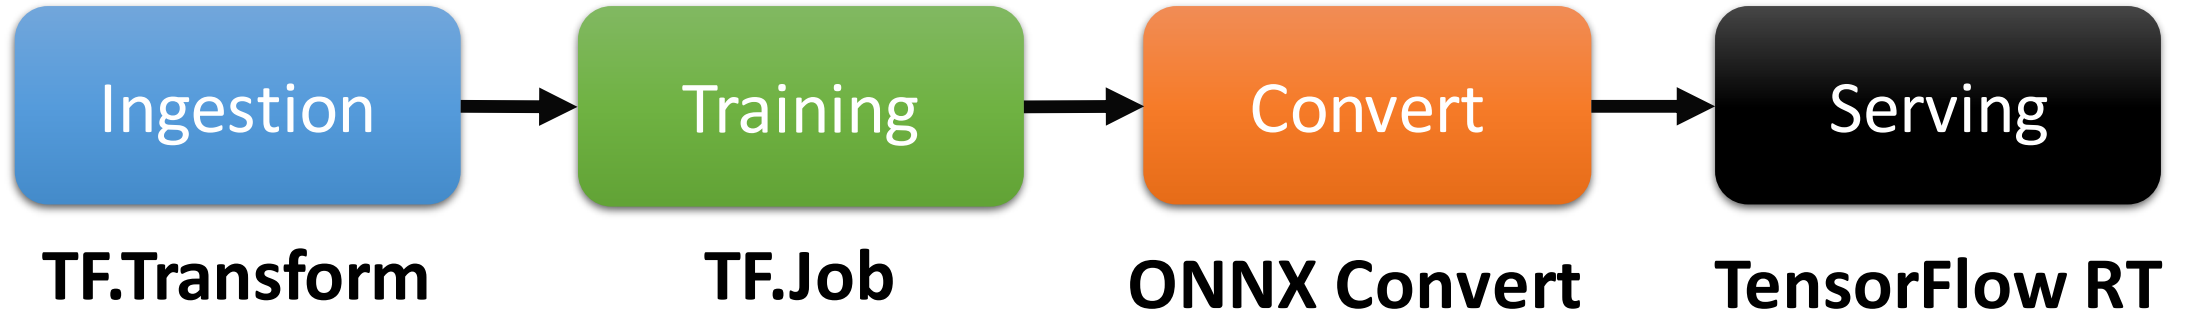
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
servingStep = dsl.ContainerOp(image=trt_image, <params>,  
                               arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



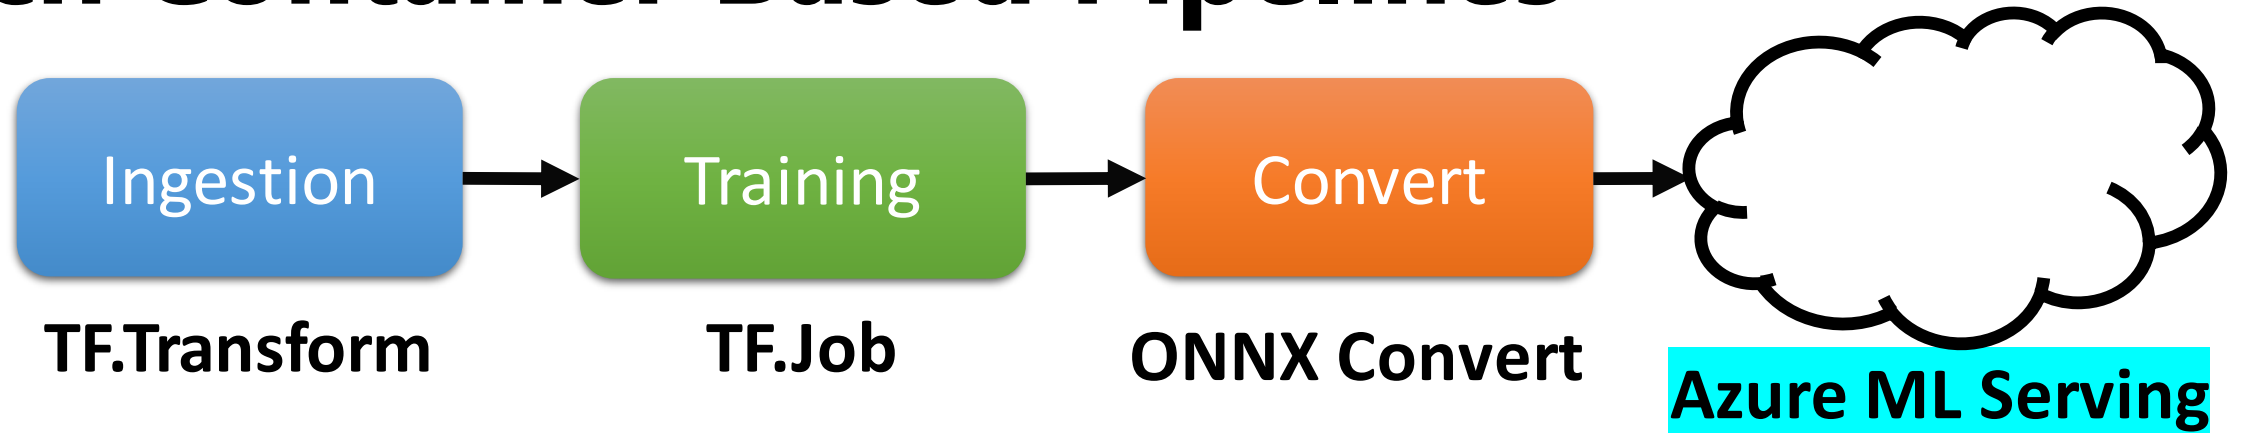
```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
deployStep = dsl.ContainerOp(image=aml_deploy_image, <params>,  
                              arguments=[convertStep.outputs['bucket']])
```

Rich Container Based Pipelines



```
ingestStep = dsl.ContainerOp(image=tft_image, <params>,  
                             file_outputs={'bucket': '/output.txt'})
```

```
trainStep = dsl.ContainerOp(image=tfjob_image, <params>,  
                             arguments=[ingestStep.outputs['bucket']])
```

```
convertStep = dsl.ContainerOp(image=convert_image, <params>,  
                               arguments=[trainStep.outputs['bucket']])
```

```
deployStep = dsl.ContainerOp(image=aml_deploy_image, <params>,  
                              arguments=[convertStep.outputs['bucket']])
```

DEMO

We're just getting started!

Our roadmap:

- Transition off of ksonnet
- Infrastructure request/provisioning via Fairing
- Improvements in the notebook manager
- **You tell us!** (Or better yet, help!)

It's a whole new world

- Data science will touch **EVERY** industry.
- We can't ask people to become a PhD in statistics though.
- How do **WE** help everyone take advantage of this transformation?



Kubeflow is **open**!



**Open
comm-
unity**



**Open
design**



**Open
to ideas**



**Open
source**

Come Help!

- website: <https://kubeflow.org>
- github: <https://github.com/kubeflow/kubeflow>
- slack: kubeflow (<http://kubeflow.slack.com>)
- twitter: @kubeflow

David Aronchick @aronchick (david.aronchick@microsoft.com)

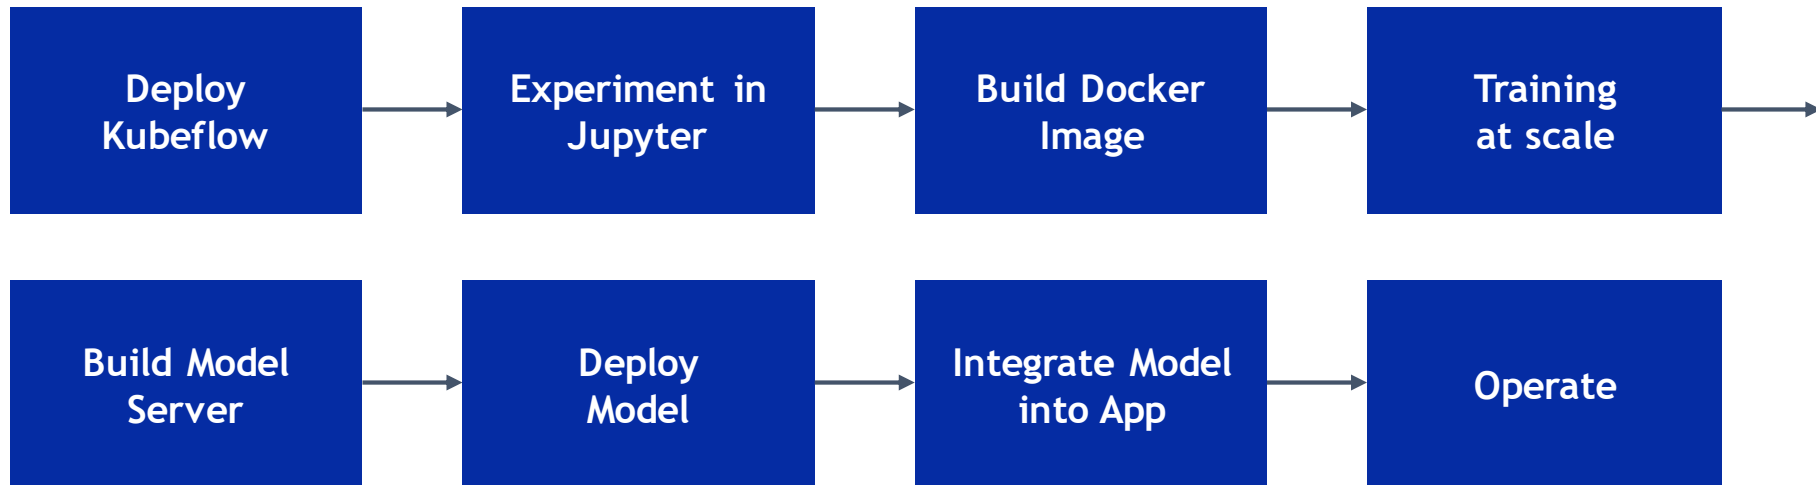
Seth Juarez (sejuarez@microsoft.com)

BONEYARD

Click to Deploy

- **Problem:** It's too hard to install Kubeflow!
- **Solution:** A one-click installation tool, available via a clean web interface
- **How:**
 - Click to deploy uses a bootstrap container and kfctl.sh with all the necessary dependencies included
 - Also enables use of declarative infrastructure deployment (e.g. Deployment Manager on GCP)
 - **NO TEMPLATING TOOL NEEDED**

User Experience



Experimentation

Model

UX

Tooling

Framework

Storage

Runtime

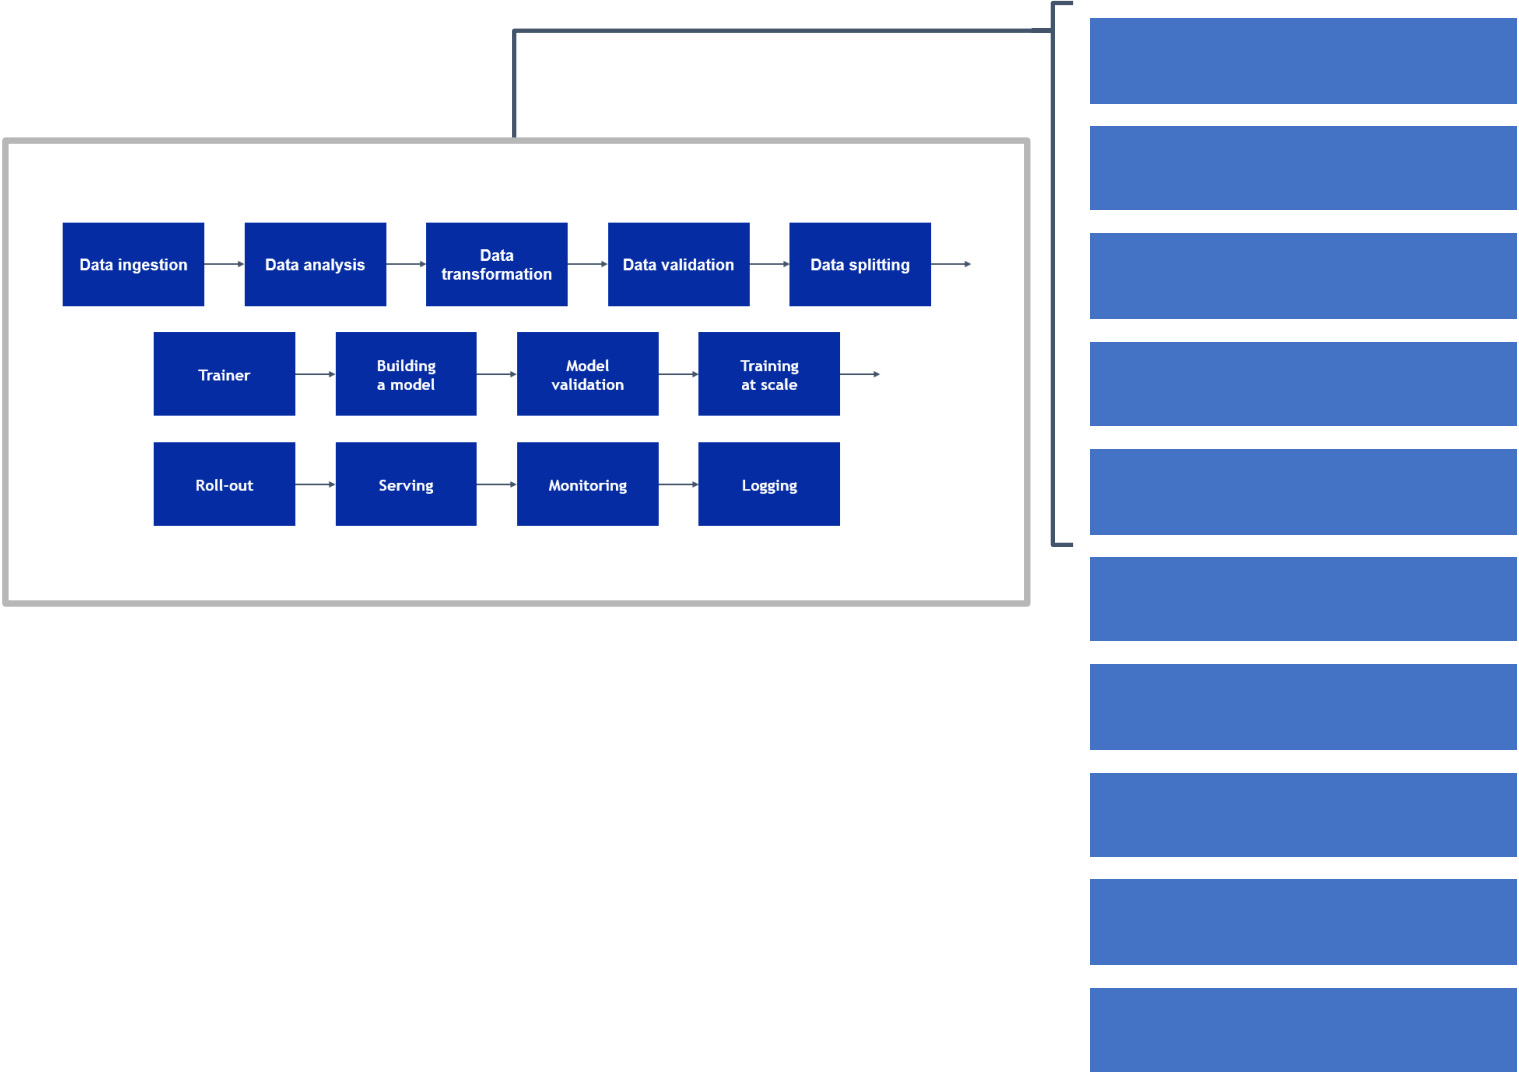
Drivers

OS

Accelerator

HW

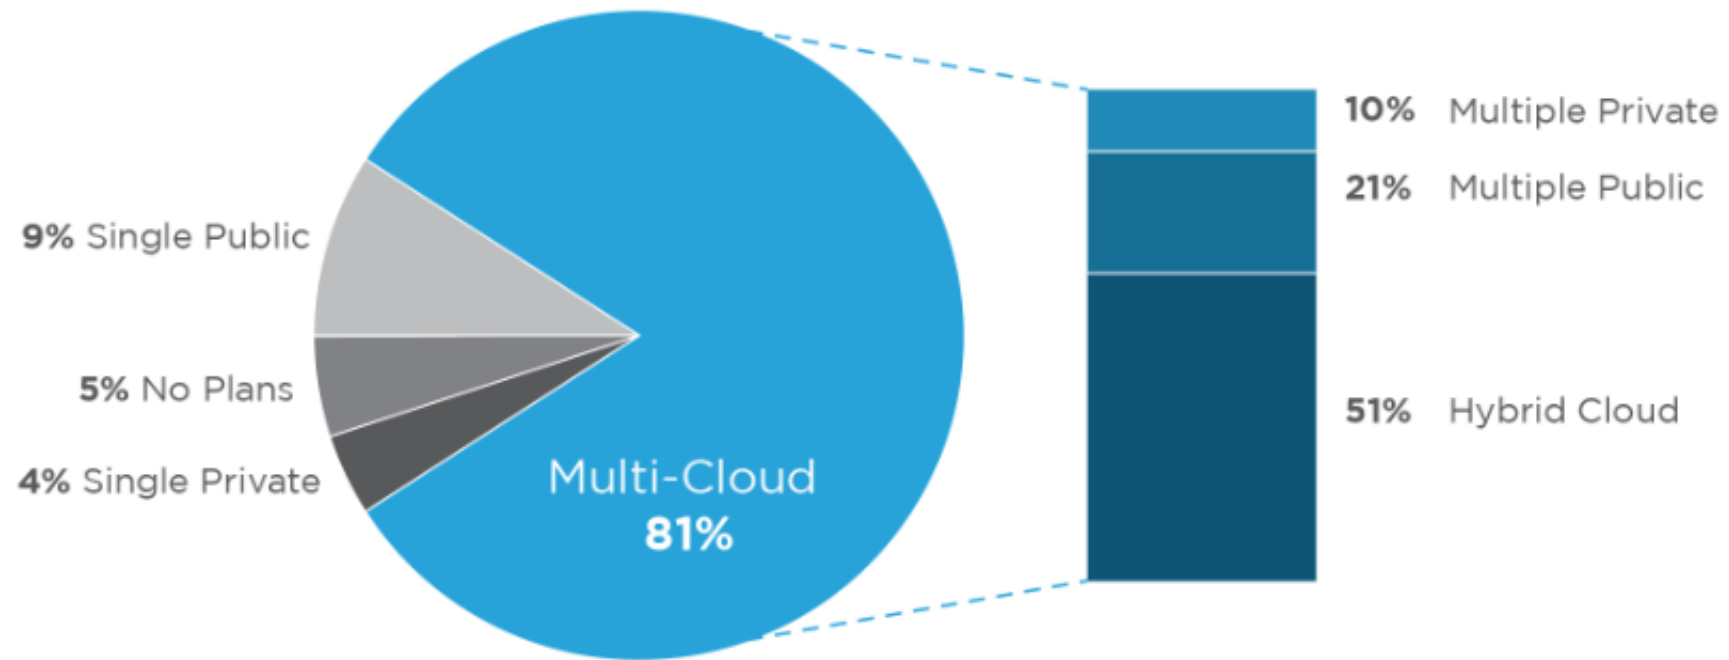
Experimentation



Multi-Cloud is the Reality

Respondents with 1,000+ Employees

81% of enterprises have a multi-cloud strategy



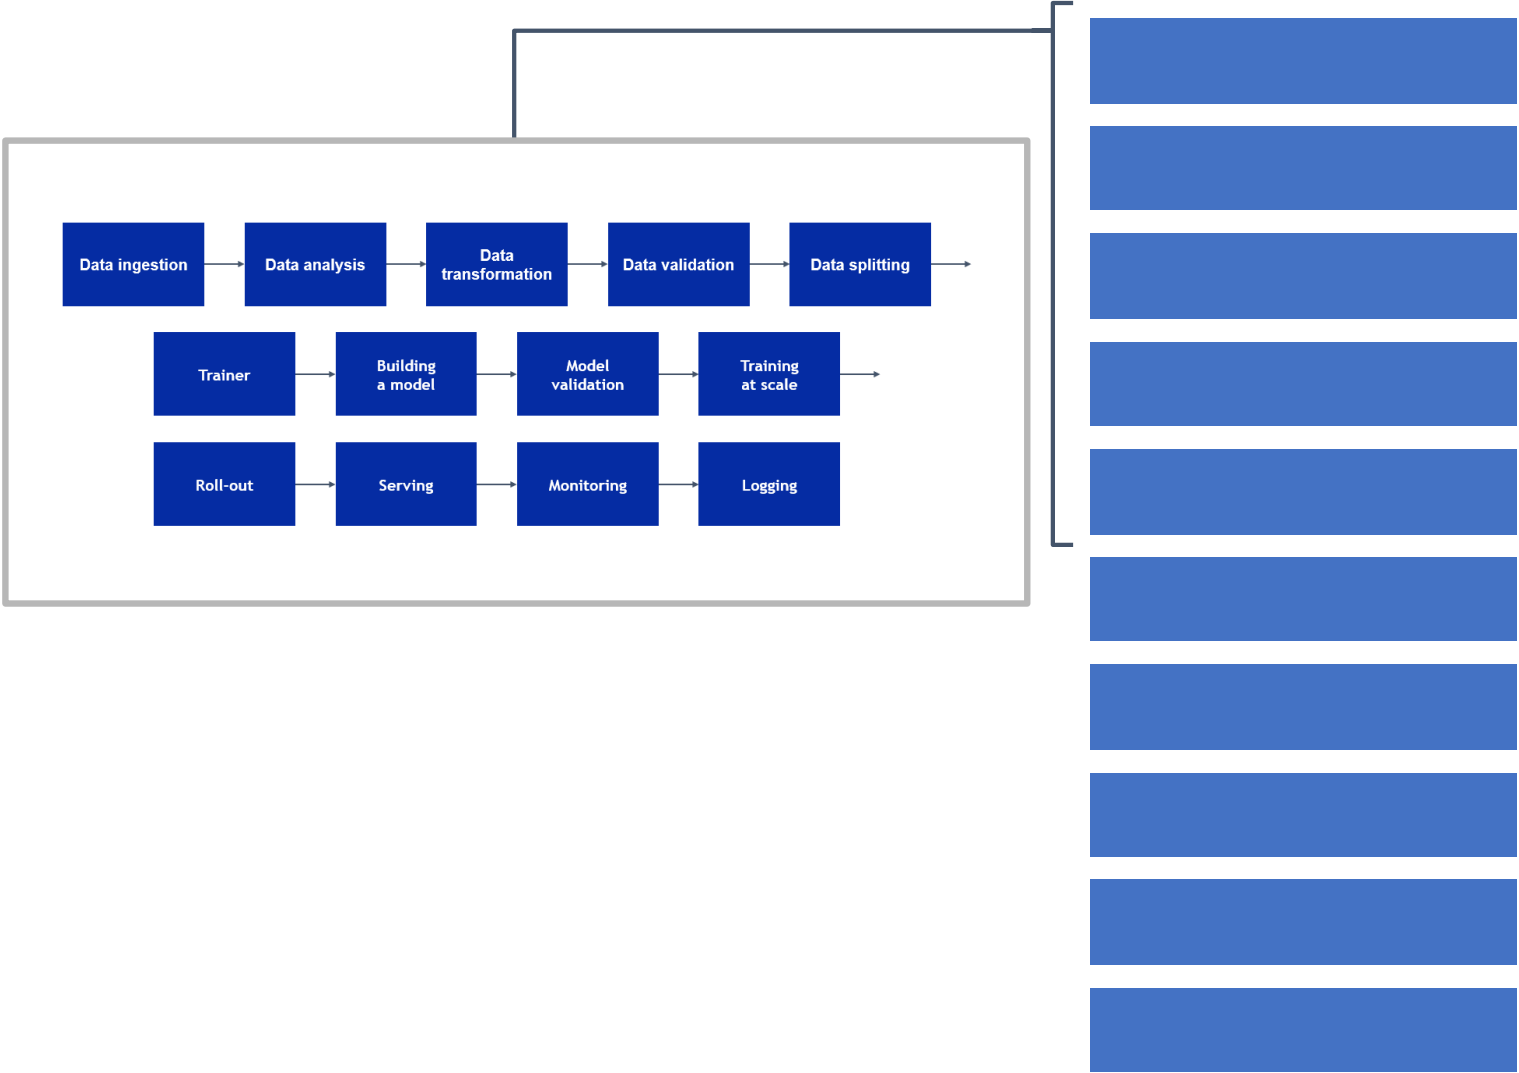
Source: RightScale 2018 State of the Cloud Report

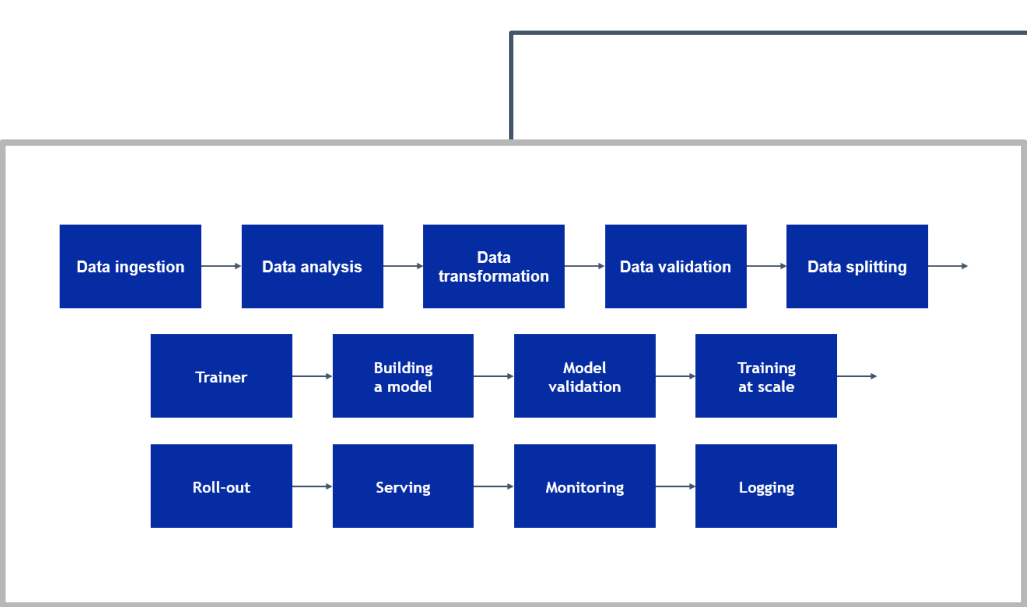
And Not Just One Cloud!

Companies using almost **5** public
and private clouds on average

Public + Private Clouds Used	Average <i>All respondents</i>	Median <i>All respondents</i>
Running Applications	3.1	3.0
Experimenting	1.7	1.0
Total	4.8	4.0

Experimentation





Experimentation

Training

